

Open Source Frameworks (OSF)

Web Services

Open Source Frameworks (OSF)
Master of Science in Engineering (MSE)
Olivier Liechti
olivier.liechti@heig-vd.ch



MASTER OF SCIENCE
IN ENGINEERING

Planning

| Date | Java EE Frameworks | Gamification Project |
|----------|---|----------------------------------|
| 23.09.13 | Intro, Java EE Overview, EJBs | Environment setup 1 |
| 30.09.13 | REST APIs & JAX-RS | Environment setup 2 (automation) |
| 07.10.13 | Design and document a REST API for your gamification engine | |
| 14.10.13 | Persistence with JPA | Test and implement your REST API |
| 21.10.13 | Break | |
| 28.10.13 | Test and implement your REST API | |
| 04.11.13 | Spring Framework | Presentations & demos |
| 11.11.13 | Technical POC Project: Define the scope & plan the activities | |
| 18.11.13 | Technical POC Project: Build the reference system | |
| 25.11.13 | Technical POC Project: Build the test infrastructure | |
| 02.12.13 | Technical POC Project: Present the results (with a demo) | |
| 09.12.13 | Introduction to Javascript frameworks | Get ready with node.js & express |
| 16.12.13 | Re-implement your REST API in Javascript | |
| 23.12.13 | Break | |
| 30.12.13 | | |
| 06.01.14 | Re-implement your REST API in Javascript | |
| 13.01.14 | Java Message Service | Presentations & demos |

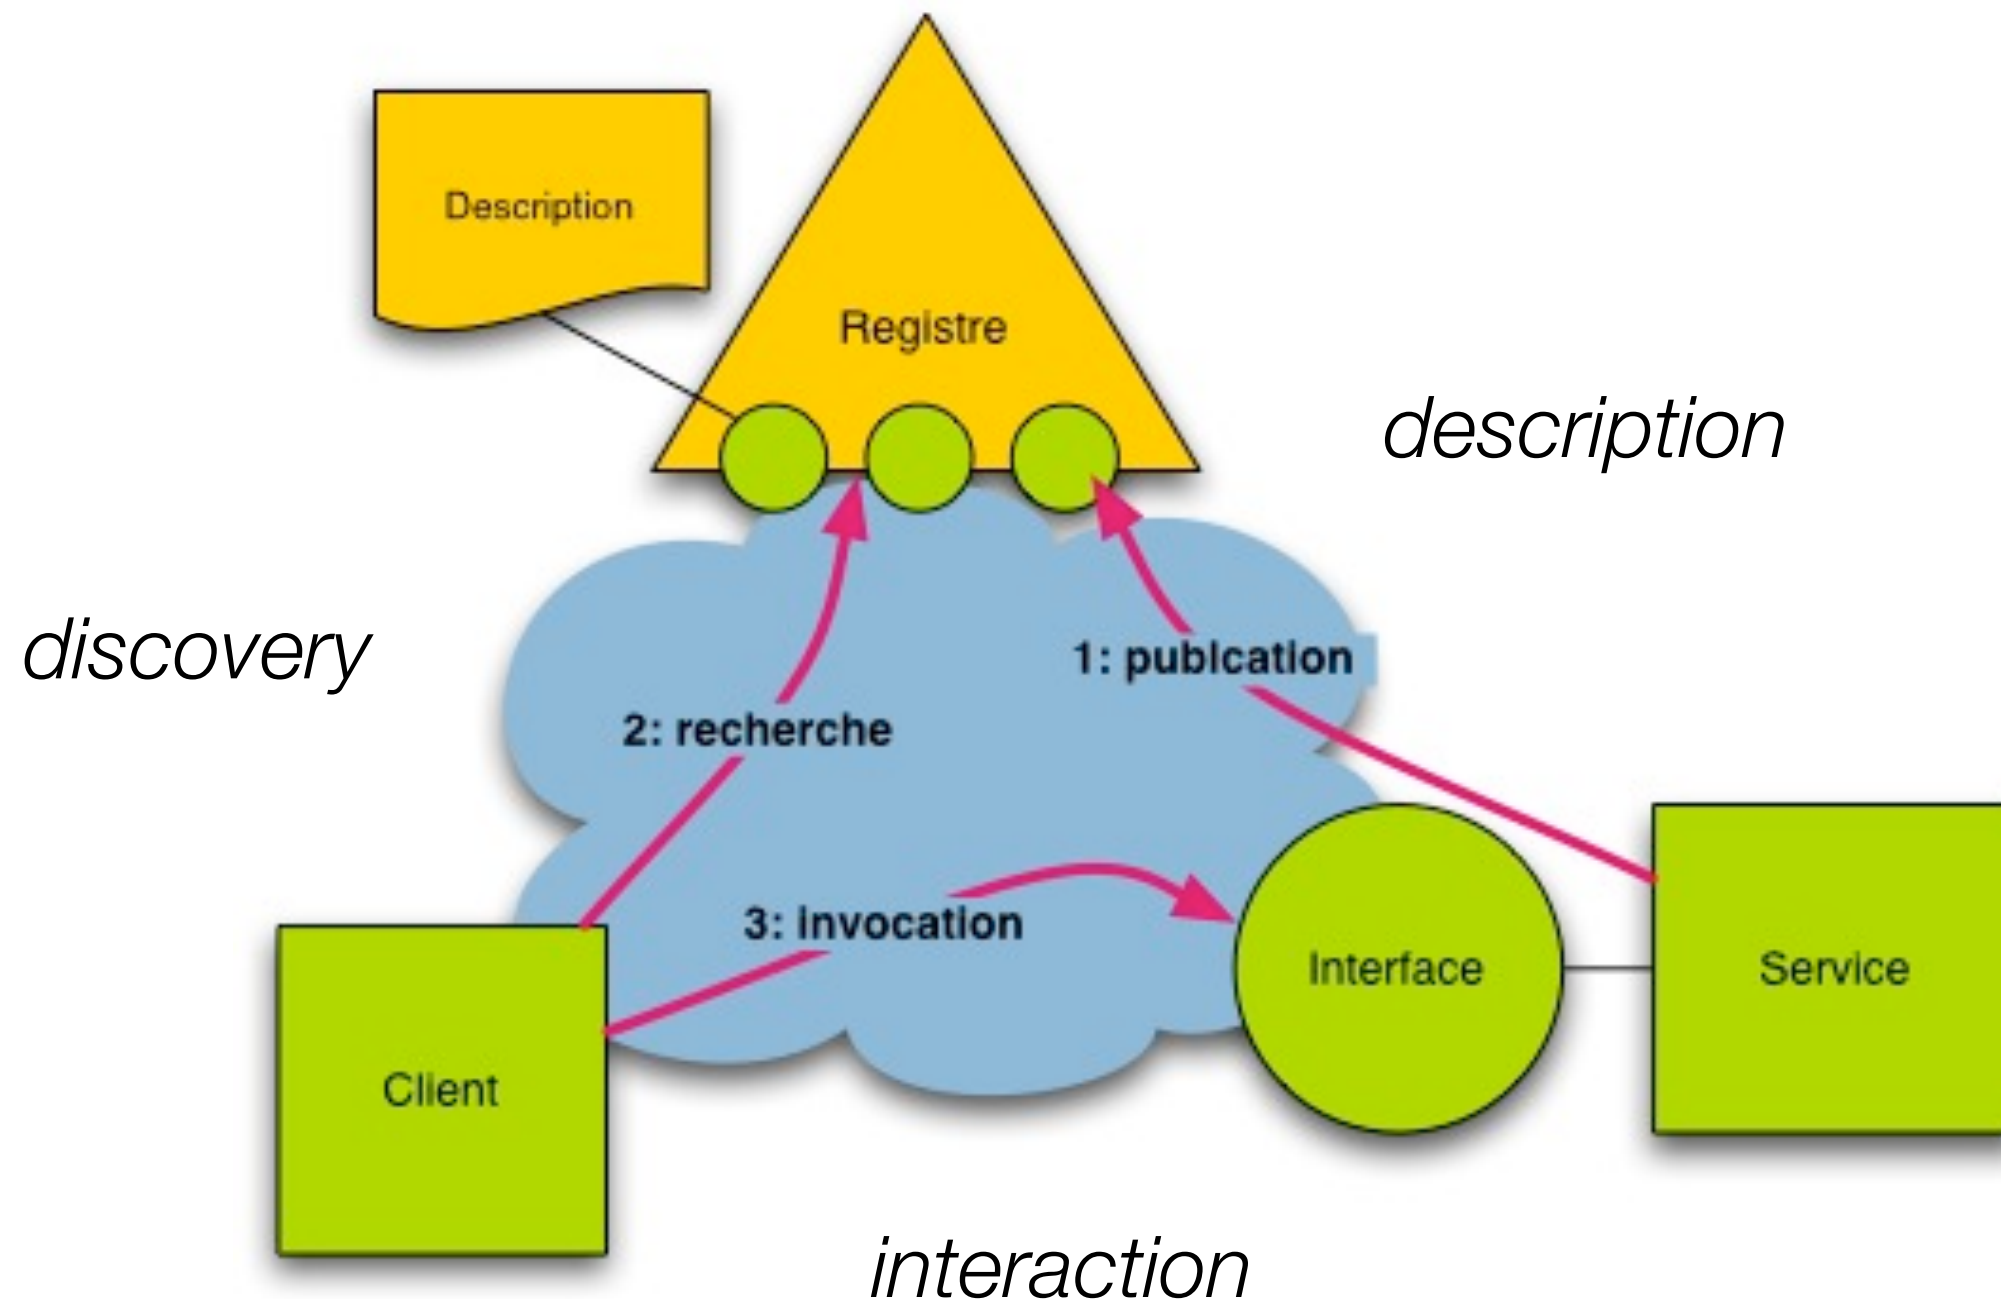
Two Approaches to Web Services



The Big Web Services Approach



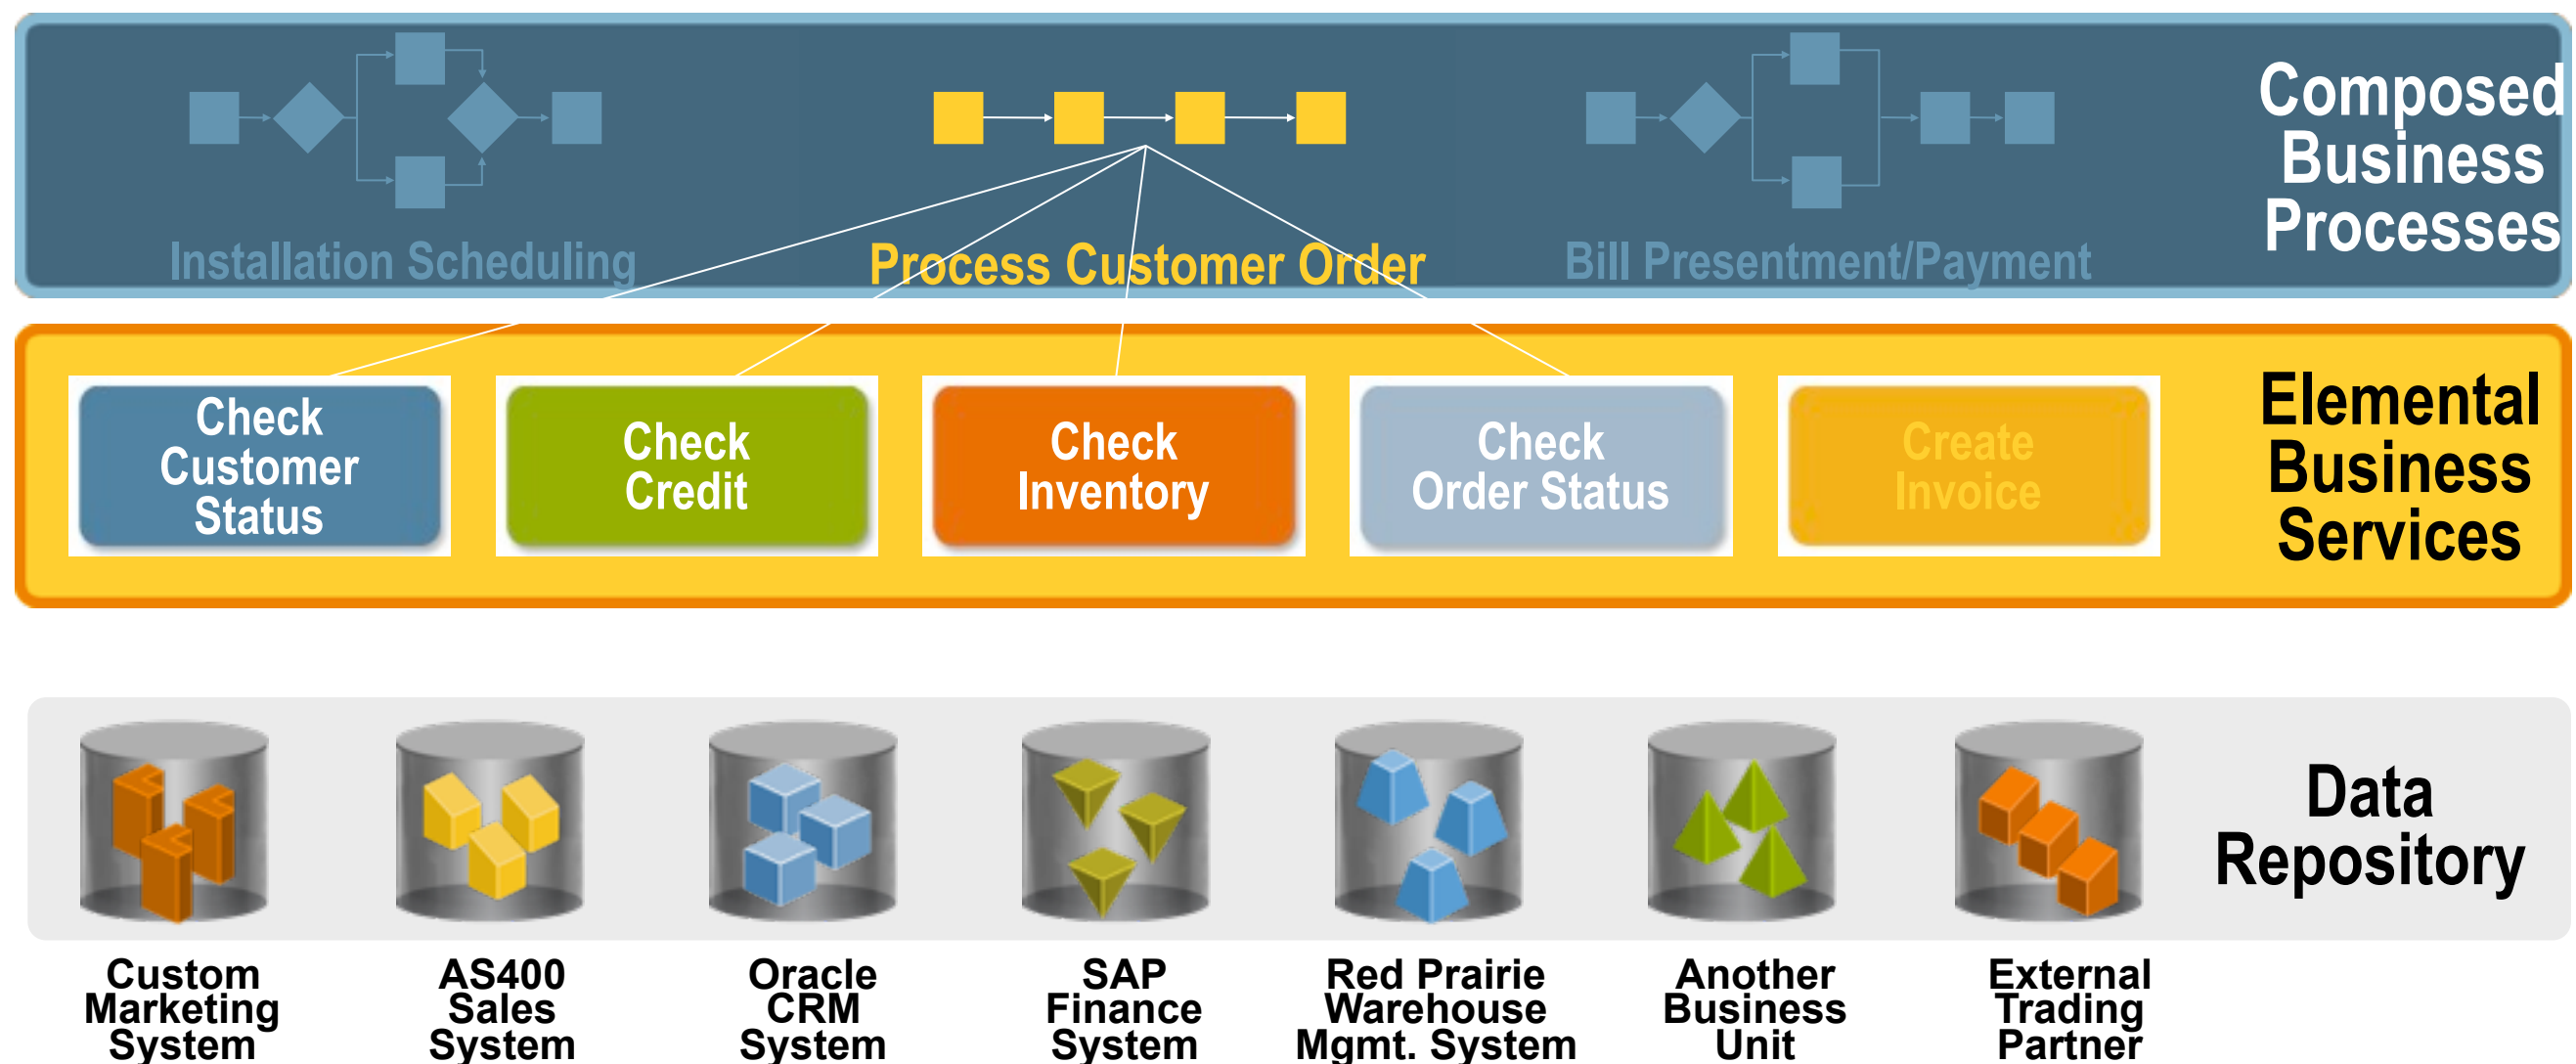

The Web Services Reference Architecture



For a **Full** Service Architecture, We Need...

- A standardized format to **describe** service interfaces
 - Example: WSDL
- A standardized protocol to **invoke** services
 - Example: SOAP
- A **registry** service
 - Example: UDDI

SOA: Service Composition & Workflows



Big Web Services with Java EE

- **JAX-WS**

- JAX-WS makes it easier to write both **web services** and **web services clients**.
- The JAX-WS **runtime** takes care of the SOAP and WSDL details and provides you with an object-oriented interface.

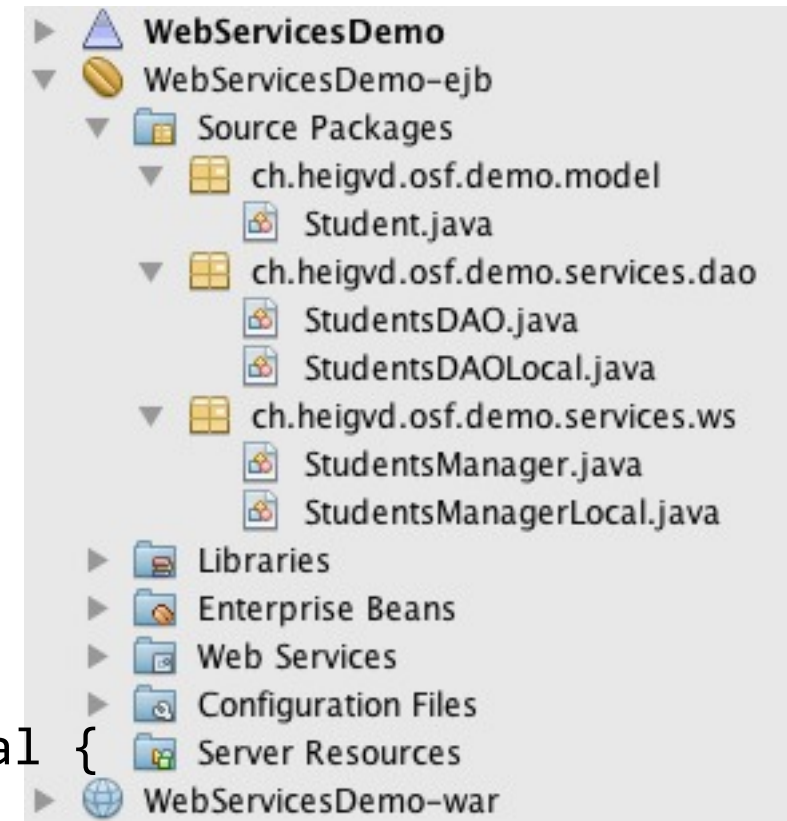
- Exposing your **Stateless Session Beans** with a Web Services interface\$

- Adding a single annotation will do the job.
- JAX-WS relies on conventions for generating the WSDL interface; you can customize the schema with various annotations.

Demo

```
@Stateless  
@WebService
```


```
public class StudentsManager implements StudentsManagerLocal {  
  
    @EJB StudentsDAOLocal studentsDAO;  
  
    public void createStudent(String firstName, String lastName) {  
        studentsDAO.createStudent(firstName, lastName);  
    }  
  
    public Student findStudentById(long id) {  
        return studentsDAO.findStudentById(id);  
    }  
}
```



The RESTful Approach



RESTful Web Services



The REST Architectural Style

- REST: REpresentational State Transfer
- REST is an architectural style for building distributed systems.
- REST has been introduced in Roy Fielding's Ph.D. thesis (Roy Fielding has been a contributor to the HTTP specification, to the apache server, to the apache community).
- The WWW is one example for a distributed system that exhibits the characteristics of a REST architecture.

HTTP is a protocol for interacting with "**resources**"

What is a “Resource”

- At first glance, one could think that a “resource” is a file on a web server:
 - an HTML document, an XML document, a PNG document
- That fits the vision of the “static content” web
- But of course, the web is now more than a huge library of hypermedia documents:
 - through the web, we interact with services and a lot of the content is dynamic.
 - more and more, through the web we interact with physical objects (machines, sensors, actuators)
 - We need a more generic definition for resources!

What is a “Resource”?

- A resource is "something" that can be named and uniquely identified:
 - Example 1: an article published in the "24 heures" newspaper
 - Example 2: the collection of articles published in the sport section of the newspaper
 - Example 3: a person's resume
 - Example 4: the current price of the Nestlé stock quote
 - Example 5: the vending machine in the school hallway
 - Example 6: the list of grades of the student Jean Dupont
- URL (Uniform Resource Locator) is a mechanism for identifying resources
 - Exemple 1: <http://www.24heures.ch/vaud/vaud/2008/08/04/trente-etudiants-partent-rencontre-patrons>
 - Exemple 2: <http://www.24heures.ch/articles/sport>
 - Exemple 5: <http://www.smart-machines.ch/customers/heig/machines/8272>

Resource vs. Representation

- A "resource" can be something intangible (stock quote) or tangible (vending machine)
- The HTTP protocol supports the exchange of data between a client and a server.
- Hence, what is exchanged between a client and a server is **not** the resource. It is a **representation** of a resource.
- Different representations of the same resource can be generated:
 - HTML representation
 - XML representation
 - PNG representation
 - WAV representation
- **HTTP provides the content negotiation mechanisms!!**

How Do We Interact With Resources?

- The HTTP protocol defines the standard methods. These methods enable the interactions with the resources:
 - **GET**: retrieve whatever information is identified by the Request-URI
 - **POST**: used to request that the origin server accept the entity enclosed in the request as a new subordinate of the resource identified by the Request-URI in the Request-Line
 - **PUT**: requests that the enclosed entity be stored under the supplied Request-URI.
 - **DELETE**: requests that the origin server delete the resource identified by the Request-URI.
 - **HEAD**: identical to GET except that the server MUST NOT return a message-body in the response
 - **TRACE**: used for debugging (echo)
 - **CONNECT**: reserved for tunneling purposes

Principles of a REST Architecture

- The state of the application is captured in a **set of resources**
 - Users, photos, comments, tags, albums, etc.
- Every resource can be **identified with a standard format** (e.g. URL)
- Every resource can have **several representations**
- There is one **unique interface for interacting** with resources (e.g. HTTP methods)
- The communication protocol is:
 - client-server
 - stateless
 - cacheable
- These properties have a positive impact on systemic qualities (scalability, performance, availability, etc.).
 - Reference: http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm

Reference

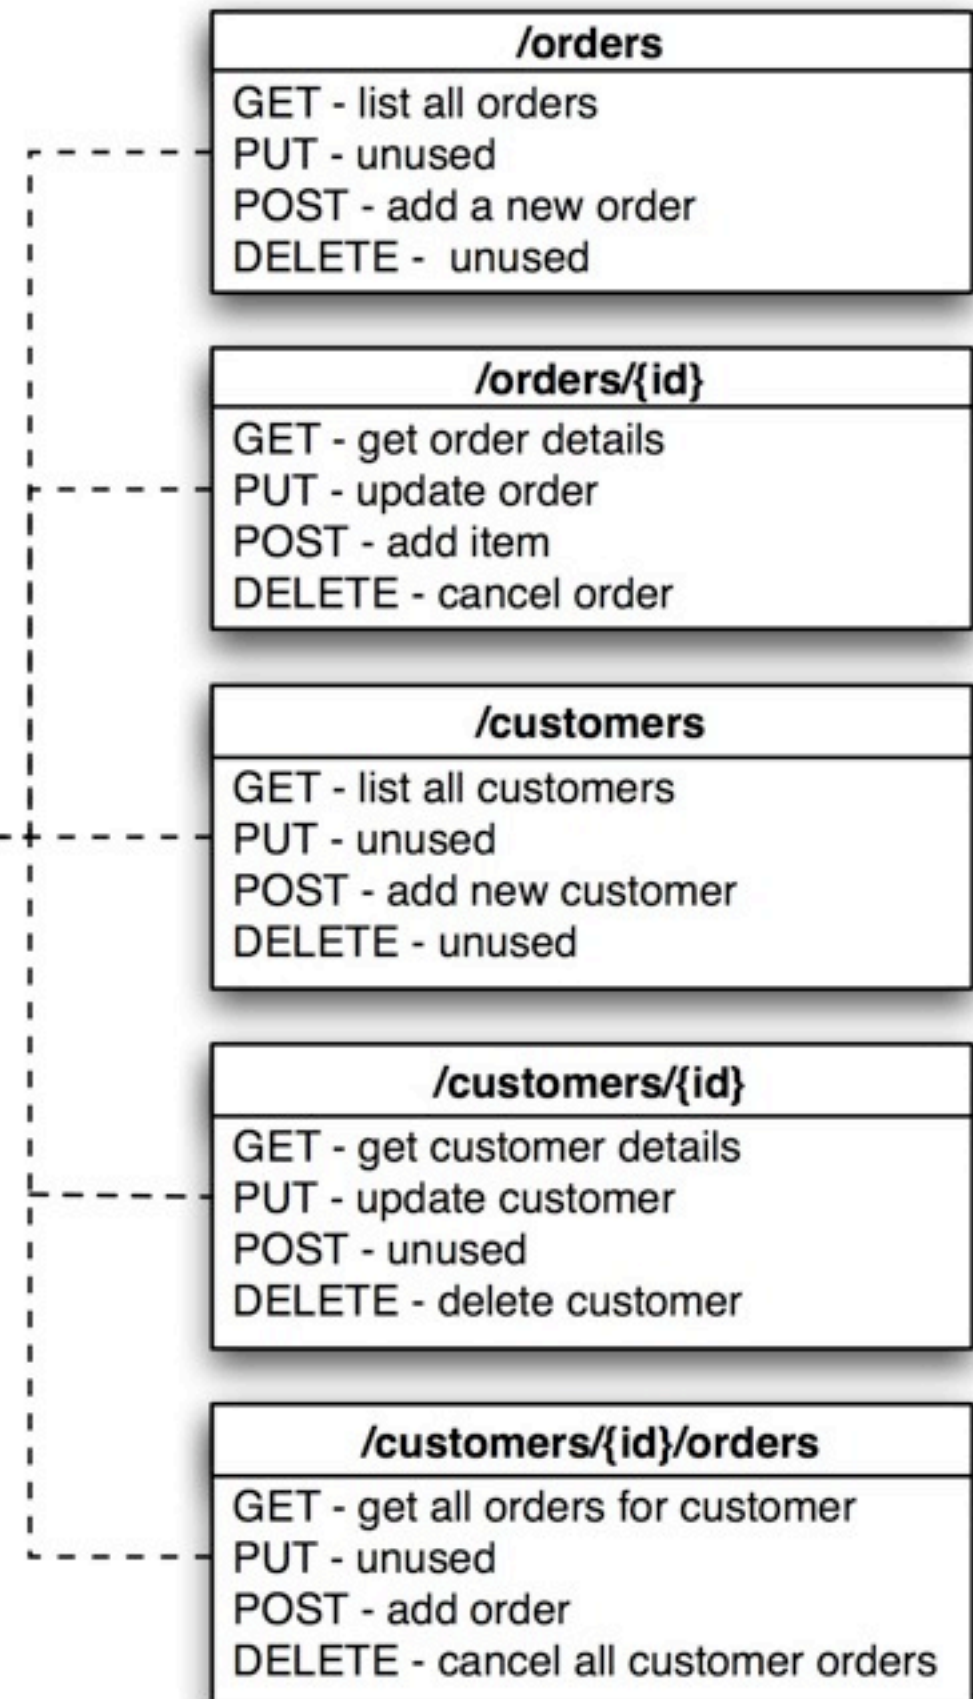

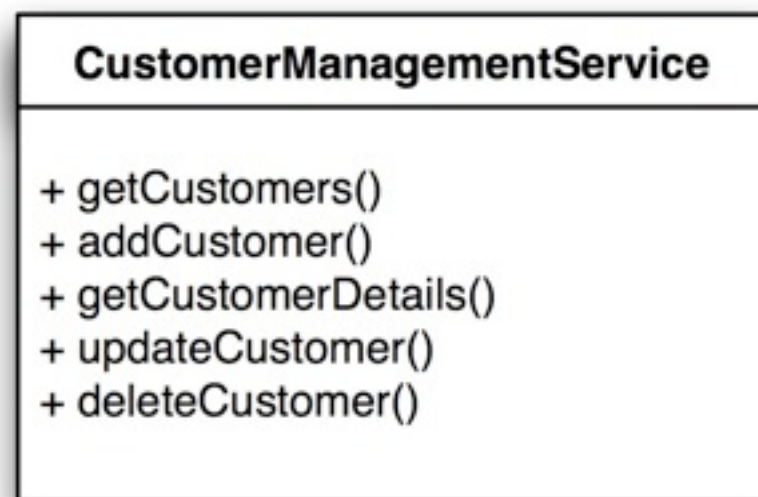
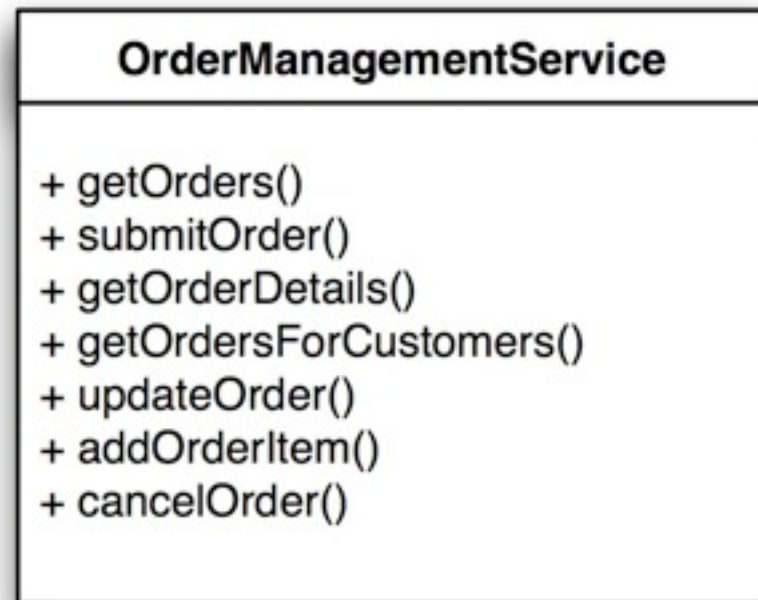
- Very good article, with presentation of key concepts and illustrative examples:
 - <http://www.infoq.com/articles/rest-introduction>

How should I specify/document my REST API?


Design a RESTful system

- Start by identifying the resources - the NAMES in your system.
- Define the structure of the URLs that will be mapped to your resources.
- Define the semantic of the operations that you want to support on all of your resources (you don't want to support GET, POST, PUT, DELETE on all resources!).
- Some examples:
 - `http://www.photos.com/users/oliehti` identifies a resource of type "user". A client can do a "HTTP GET" to obtain a representation of the user or a "HTTP PUT" to update the user.
 - `http://www.photos.com/users` identifies a resource of type "collection of users". A client can do a "HTTP POST" to add users, or an "HTTP GET" to obtain the list of users.

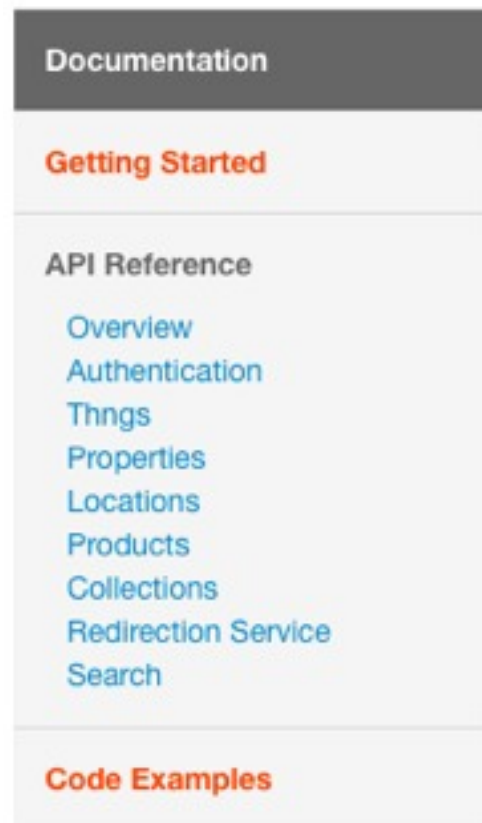
RPC vs REST



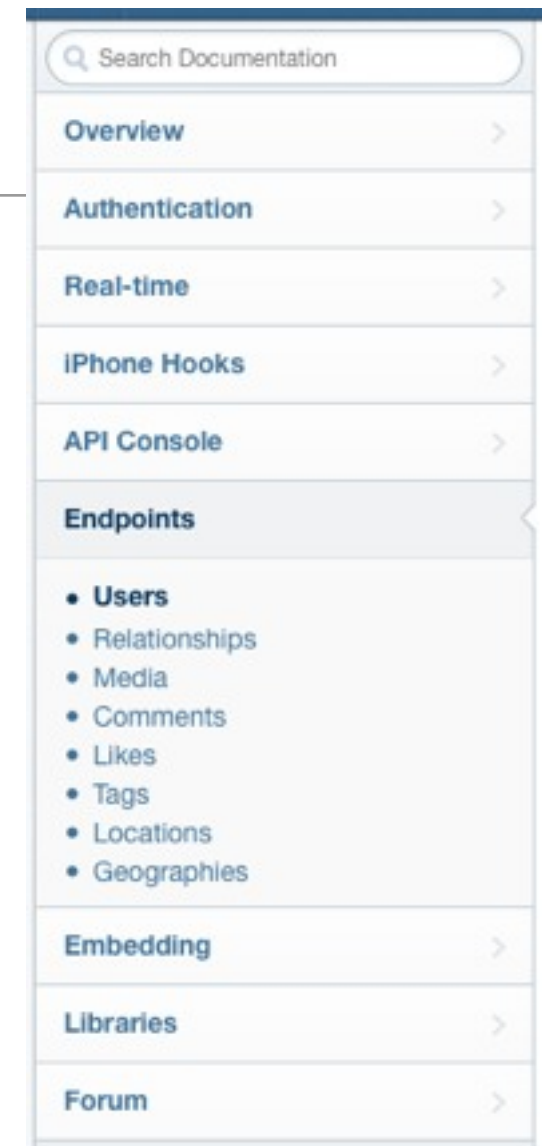
Look at Some Examples



<http://dev.librato.com/v1>



<https://dev.evrythng.com/documentation/api>



<http://instagram.com/developer/endpoints/>



What Are Metrics?

Metrics are custom measurements stored in Librato's Metrics service. These measurements are created and may be accessed programmatically through a set of RESTful API calls. There are currently two types of metrics that may be stored in Librato Metrics, **gauges** and **counters**.

Gauges

Gauges capture a series of measurements where each measurement represents the value under observation at one point in time. The value of a gauge typically varies between some known minimum and maximum. Examples of gauge measurements include the requests/second serviced by an application, the amount of available disk space, the current value of \$AAPL, etc.

Counters

Counters track an increasing number of occurrences of some event. A counter is unbounded and always monotonically increasing in any given run. A new run is started anytime that counter is reset to zero. Examples of counter measurements include the number of connections made to an app, the number of visitors to a website, the number of a times a write operation failed, etc.

Metric Properties

Some common properties are supported across all types of metrics:

name

Each metric has a name that is unique to its class of metrics e.g. a gauge name must be unique to gauges. The name identifies a metric in subsequent API calls to store/query individual measurements. The name can be up to 63 characters in length. Valid characters for metric names are 'A-Za-z0-9.-_,'.

period

The **period** of a metric is an integer value that describes (in seconds) the standard reporting interval for the metric. Setting the period enables Metrics to detect abnormal interruptions in reporting and to

GET /v1/metrics/:name

API VERSION 1.0

Description

Returns information for a specific metric. If time interval search parameters are specified will also include a set of metric measurements for the given time span.

URL

`https://metrics-api.librato.com/v1/metrics/:name`

Method

GET

Measurement Search Parameters

If optional **time interval search parameters** are specified, the response includes the set of metric measurements covered by the time interval. Measurements are listed by their originating source name if one was specified when the measurement was created. All measurements that were created without an explicit source name are listed with the source name **unassigned**.

source

Deprecated: Use **sources** with a single source name, e.g [mysource].

sources

If **sources** is specified, the response is limited to measurements from those sources. The **sources** parameter should be specified as an array of source names. The response is limited to the set of sources specified in the array.

Examples &
payload structure

Examples

Return the metric named `cpu_temp` with up to four measurements at resolution 60.

```
curl \
  -u <user>:<token> \
  -X GET \
  https://metrics-api.librato.com/v1/metrics/cpu_temp?resolution=60&count=4
```

Response Code

200 OK

Response Headers

** NOT APPLICABLE **

Response Body

```
{
  "type": "gauge",
  "display_name": "cpu_temp",
  "resolution": 60,
  "measurements": [
    {
      "source": "librato.com": {
        "name": "cpu_temp",
        "value": 84.5,
        "time": 1234567890,
        "resolution": 60
      },
      "source": "librato.com": {
        "name": "cpu_temp",
        "value": 86.7,
        "time": 1234567950,
        "resolution": 60
      },
      "source": "librato.com": {
        "name": "cpu_temp",
        "value": 84.6,
        "time": 1234568010,
        "resolution": 60
      },
      "source": "librato.com": {
        "name": "cpu_temp",
        "value": 89.7,
        "time": 1234568070,
        "resolution": 60
      }
    ],
    "format": "json",
    "ts_short": "8#176;F",
    "ua": "librato-metrics/0.7.4 (ruby; 1.9.3p194; x86_64-linux) direct-faraday/0.8.4",
    "ts": null,
    "ts_long": "Fahrenheit",
    "checked": true
  },
  "description": "Current CPU temperature in Fahrenheit",
  "temp"
}
```

navigator

What are Metrics?

Authentication

Response Codes & Errors

Pagination

Time Intervals

Metric Attributes

- Metrics

GET /metrics

POST /metrics

DELETE /metrics

GET /metrics/iname

PUT /metrics/iname

DELETE /metrics/iname

+ Instruments

+ Dashboards

+ Tags

+ Alerts

+ Services

+ Annotations

+ Chart Tokens

+ Users



Short description of the whole domain model

Overview

The central data structure in our engine are **Things**, which are data containers to store all the data generated by and about any physical object. Various **Properties** can be attached to any Thing, and the content of each property can be updated any time, while preserving the history of those changes. Things can be added to various **Collections** which makes it easier to share a set of Things with other **Users** within the engine.

Thing

An abstract notion of an object which has location & property data associated to it. Also called Active Digital Identities (ADIs), these resources can model real-world elements such as persons, places, cars, guitars, mobile phones, etc.

Property

A Thing has various properties: arbitrary key/value pairs to store any data. The values can be updated individually at any time, and can be retrieved historically (e.g. "Give me the values of property X between 10 am and 5 pm on the 16th August 2012").

Location

Each Thing also has a special type of Properties used to store snapshots of its geographic position over time (for now only GPS coordinates - latitude and longitude).

User

Each interaction with the EVERYTHING back-end is authenticated and a user is associated with each action. This dictates security access.

Collection

A collection is a grouping of Things. Call one collection.

Creating a new Product

To create a new **Product**, simply POST a JSON document that describes a product to the **/products** endpoint.

```
POST /products
Content-Type: application/json
Authorization: $EVERYTHING_API_KEY
```

```
{
  "fn": <String>,
  "description": <String>,
  "brand": <String>,
  "categories": [<String>, ...],
  "photos": [<String>, ...],
  "url": <String>,
  "identifiers": {
    <String>: <String>,
    ... },
  "properties": {
    <String>: <String>,
    ... },
  "tags": [<String>, ...]
}
```

Mandatory Parameters

fn

<String> The functional name of the product.

Optional Parameters

description

<String> An string that gives more details about the product, a short description.

CRUD method description

More details about the Product resource (domain model) & payload structure

Products

Products are very similar to things, but instead of modeling an individual object instance, products are used to model a class of objects. Usually, they are used for general classes of things, usually a particular model with specific characteristics. Let's take for example a specific TV model (e.g. [this one](#)), which has various properties such as a model number, a description, a brand, a category, etc. Products are useful to capture the properties that are common to a set of things (so you don't replicate a property "model name" or "weight" for thousands of things that are individual instances of a same product category).

The Product document model used in our engine has been designed to be compatible with the [hProduct microformat](#), therefore it can easily be integrated with the hProduct data model and applications supporting microformats.

The Product document model is as follows:

```
<Product>={
  "id": <String>,
  "createdAt": <timestamp>,
  "updatedAt": <timestamp>,
  "fn": <String>,
  "description": <String>,
  "brand": <String>,
  "categories": [<String>, ...],
  "photos": [<String>, ...],
  "url": <String>,
  "identifiers": {
    <String>: <String>,
    ... },
  "properties": {
    <String>: <String>,
    ... },
  "tags": [<String>, ...]
}
```

Cross-cutting concerns

Pagination

Requests that return multiple items will be paginated to 30 items by default. You can specify further pages with the **?page** parameter. You can also set a custom page size up to 100 with the **?per_page** parameter.

Authentication

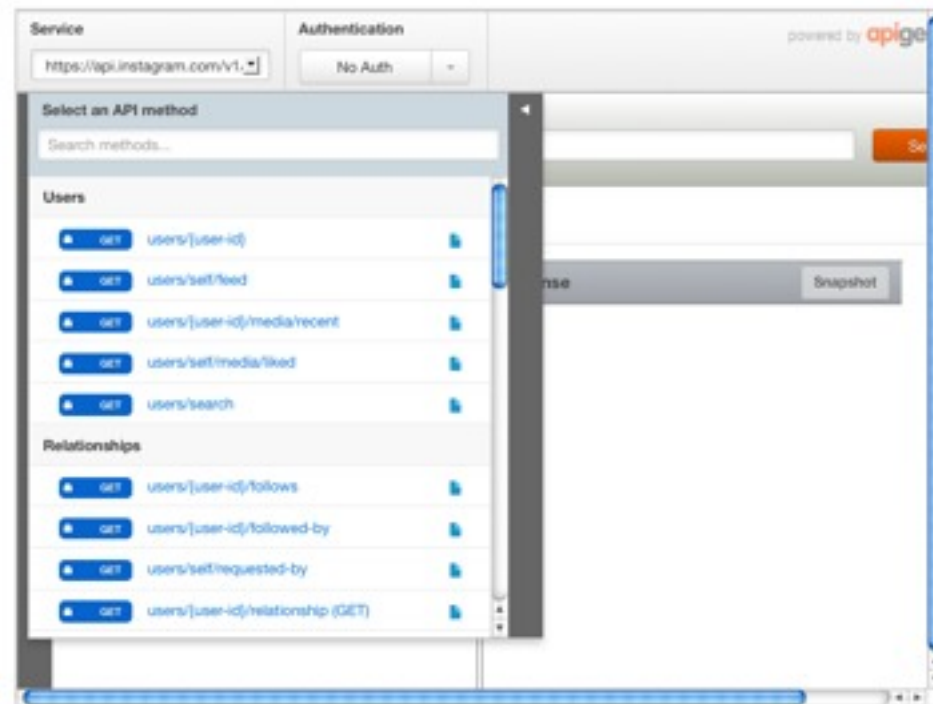
Access to our API is done via HTTPS requests to the <https://api.everything.com> domain. Unencrypted HTTP requests are accepted (<http://api.everything.com> for low-power device without SSL support), but we **strongly** suggest to use only HTTPS if you store any valuable data in our engine. Every request to our API must include an API key using **Authorization** HTTP header to identify the user or application issuing the request and execute it if authorized.



Interactive test console

API Console

Our API console is provided by Apigee. Tap the Lock icon, select OAuth, and you can experiment with making requests to our API. [See it in full screen](#) →



List of supported CRUD methods for each resource (R, R/W)

User Endpoints

| | | |
|-----|--------------------------------------|---|
| GET | /users/ user-id | ... Get basic information about a user. |
| GET | /users/self/feed | ... See the authenticated user's feed. |
| GET | /users/ user-id /media/recent | ... Get the most recent media published by a user. |
| GET | /users/self/media/liked | ... See the authenticated user's list of liked media. |
| GET | /users/search | ... Search for a user by name. |

Comment Endpoints

| | | |
|------|--|--|
| GET | /media/ media-id /comments | ... Get a full list of comments on a media. |
| POST | /media/ media-id /comments | ... Create a comment on a media. Please email apide... |
| DEL | /media/ media-id /comments/ comment-id | ... Remove a comment. |

GET /media/ **media-id** /comments

https://api.instagram.com/v1/media/555/comments?access_token=ACCESS-TOKEN

RESPONSE

```
{
  "meta": {
    "code": 200
  },
  "data": [
    {
      "created_time": "1288788324",
      "text": "Really amazing photo!",
      "from": {
        "username": "snoopdogg",
        "profile_picture": "http://images.instagram.com/profiles/profile_16_75sq_1385612434.jpg",
        "id": "1574883",
        "full_name": "Snoop Dogg"
      },
      "id": "420"
    },
    ...
  ]
}
```

Get a full list of comments on a media.
Required scope: comments

CRUD method description

Cross-cutting concerns

Limits

Be nice. If you're sending too many requests too quickly, we'll send back a 503 error code (server unavailable).

You are limited to 5000 requests per hour per access_token or client_id overall. Practically, this means you should (when possible) authenticate users so that limits are well outside the reach of a given user.

PAGINATION


Sometimes you just can't get enough. For this reason, we've provided a convenient way to access more data in any request for sequential data. Simply call the url in the next_url parameter and we'll respond with the next set of data.

The Envelope

Every response is contained by an envelope. That is, each response has a predictable set of keys with which you can expect to interact:

```
{
  "meta": {
    "code": 200
  },
  "data": {
    ...
  },
  "pagination": {
    "next_url": "...",
    "next_max_id": "13872296"
  }
}
```

Some Tools that Might Help/Inspire You



<http://apigee.com/docs/>

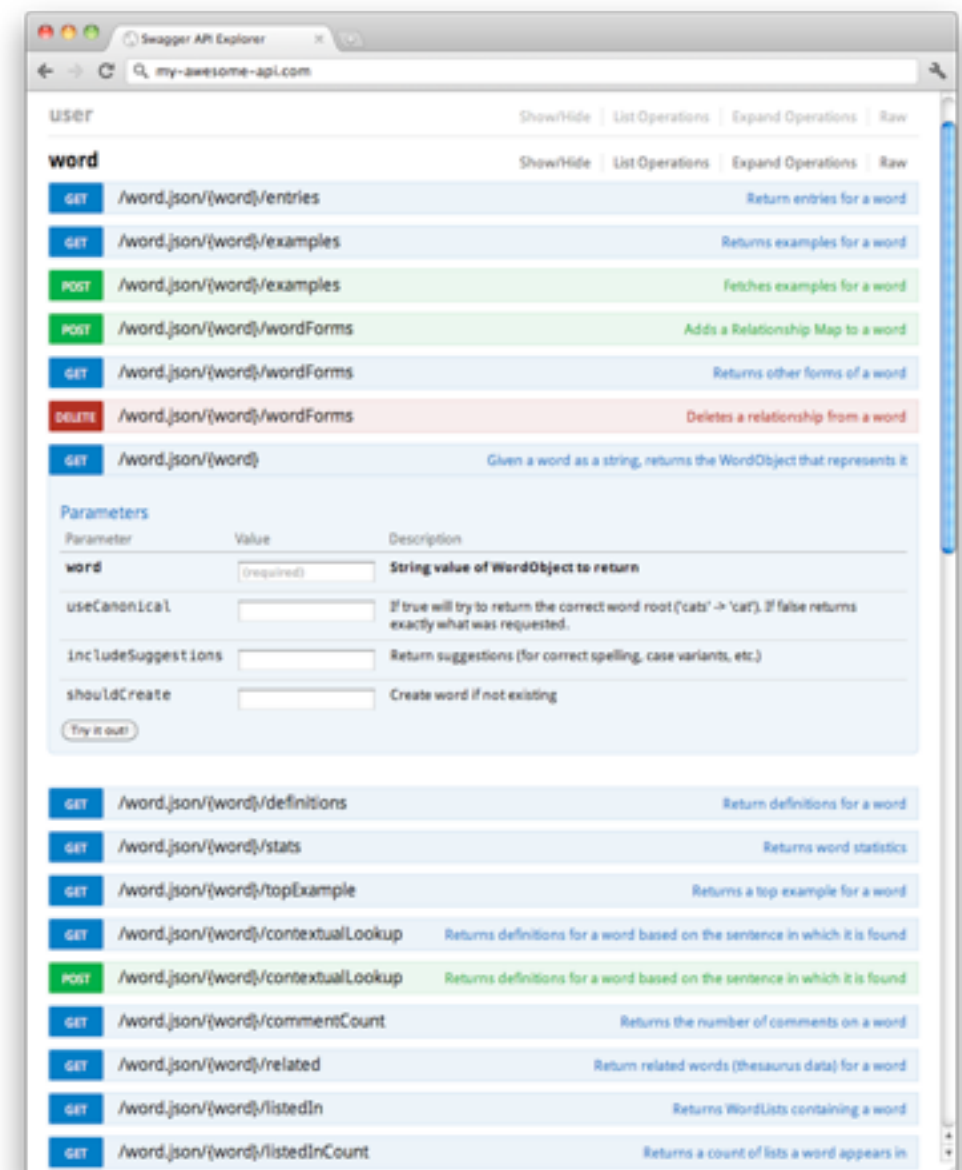


REST API documentation. Reimagined.

It takes more than a simple HTML page to thrill your API users. The right tools take weeks of development. Weeks that apiary.io saves.

```
GET /shopping-cart
> Accept: application/json
< 200
< Content-Type: application/json
{ "items": [
  { "url": "/shopping-cart/1", "product": "2ZPZ",
    "quantity": 1, "name": "New socks", "price": 1.25 }
  ] }
```

<http://apiary.io/>



<https://developers.helloverb.com/swagger/>

How to write a “RESTful” Web Service?

- On the server side, one could do everything in a FrontController servlet:
 - Parse URLs
 - Do a mapping between URLs and Java classes that represent resources
 - Generate the different representations of resources
 - etc.
- But of course, there are frameworks that do exactly that for us.
- It is true for nearly every platform and language, including Java.
- There is even a JSR for that: JAX-RS (JSR 311).
 - Oracle provides the reference implementation, in the Jersey project (open source).

Big Web Services with Java EE

- **JAX-RS**

- JAX-RS provides a programming model, classes and annotations for easily building RESTful APIs.
- Jersey is the name of the standard JAX-RS implementation, which is bundled with the Glassfish application server.

- **JAXB**

- You do not have to worry about the serialization of your business objects to XML or JSON. The framework will take care of (most of) the details for you.

- For all of your business “resources”, create a **JAX-RS resource class**

- Use **annotations** to **route HTTP requests to your resource class and methods** (based on target URI, HTTP method, HTTP accept header, etc.)



JavaOneSM

Developing RESTful Web Services with JAX-RS

Marc Hadley
Paul Sandoz
Sun Microsystems, Inc

Example

```
@Path("/students")
public class StudentsResource {

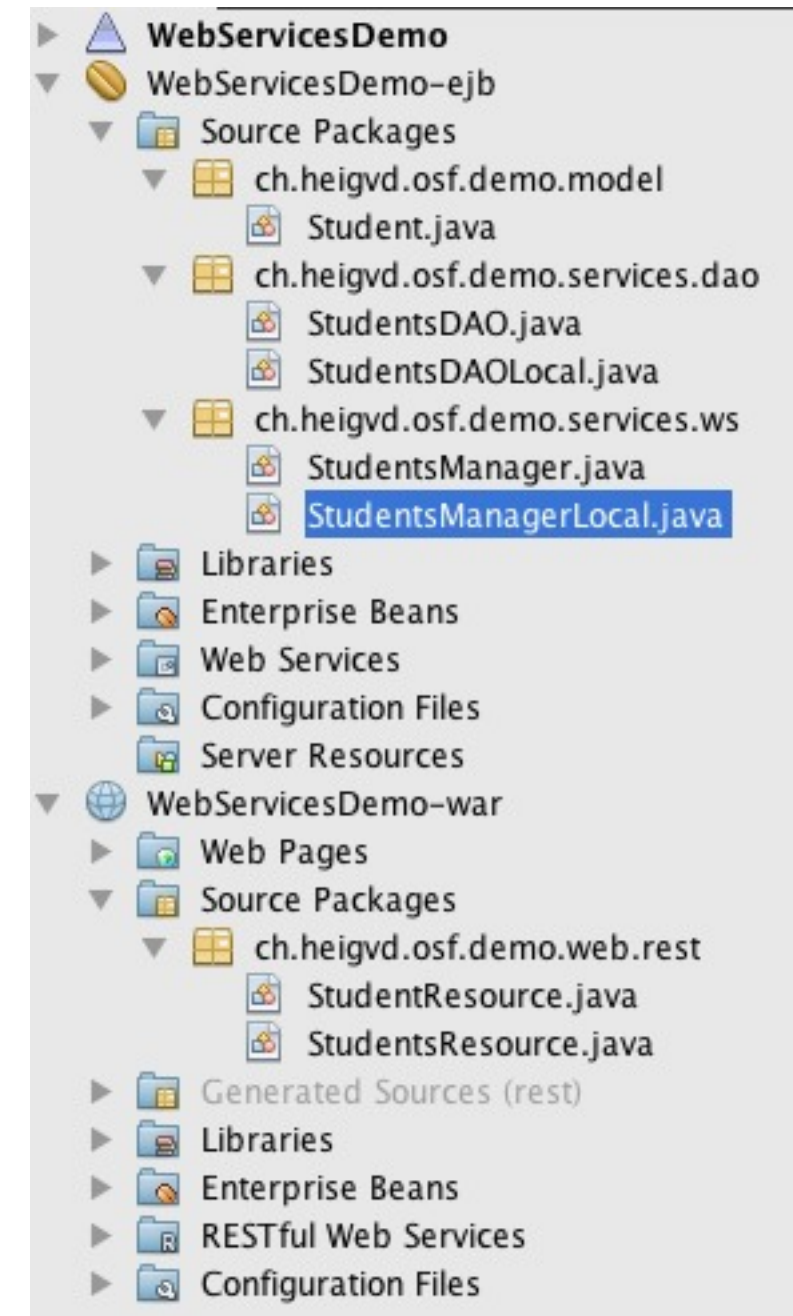
    StudentsDAOLocal studentsDAO = lookupStudentsDAOLocal();

    @Context
    private UriInfo context;


    /**
     * Creates a new instance of StudentsResource
     */
    public StudentsResource() {
    }

    /**
     * Retrieves representation of the collection resource
     * @return an instance of List<Student>
     */
    @GET
    @Produces("application/xml, application/json")
    public List<Student> getXml() {
        // Let's generate random students for demo purposes...
        // don't try to understand the logic of this
        List<Student> dummyResult = new LinkedList<Student>();
        dummyResult.add(studentsDAO.findStudentById(42));
        dummyResult.add(studentsDAO.findStudentById(42));
        dummyResult.add(studentsDAO.findStudentById(42));
        dummyResult.add(studentsDAO.findStudentById(42));
        dummyResult.add(studentsDAO.findStudentById(42));
        return dummyResult;
    }

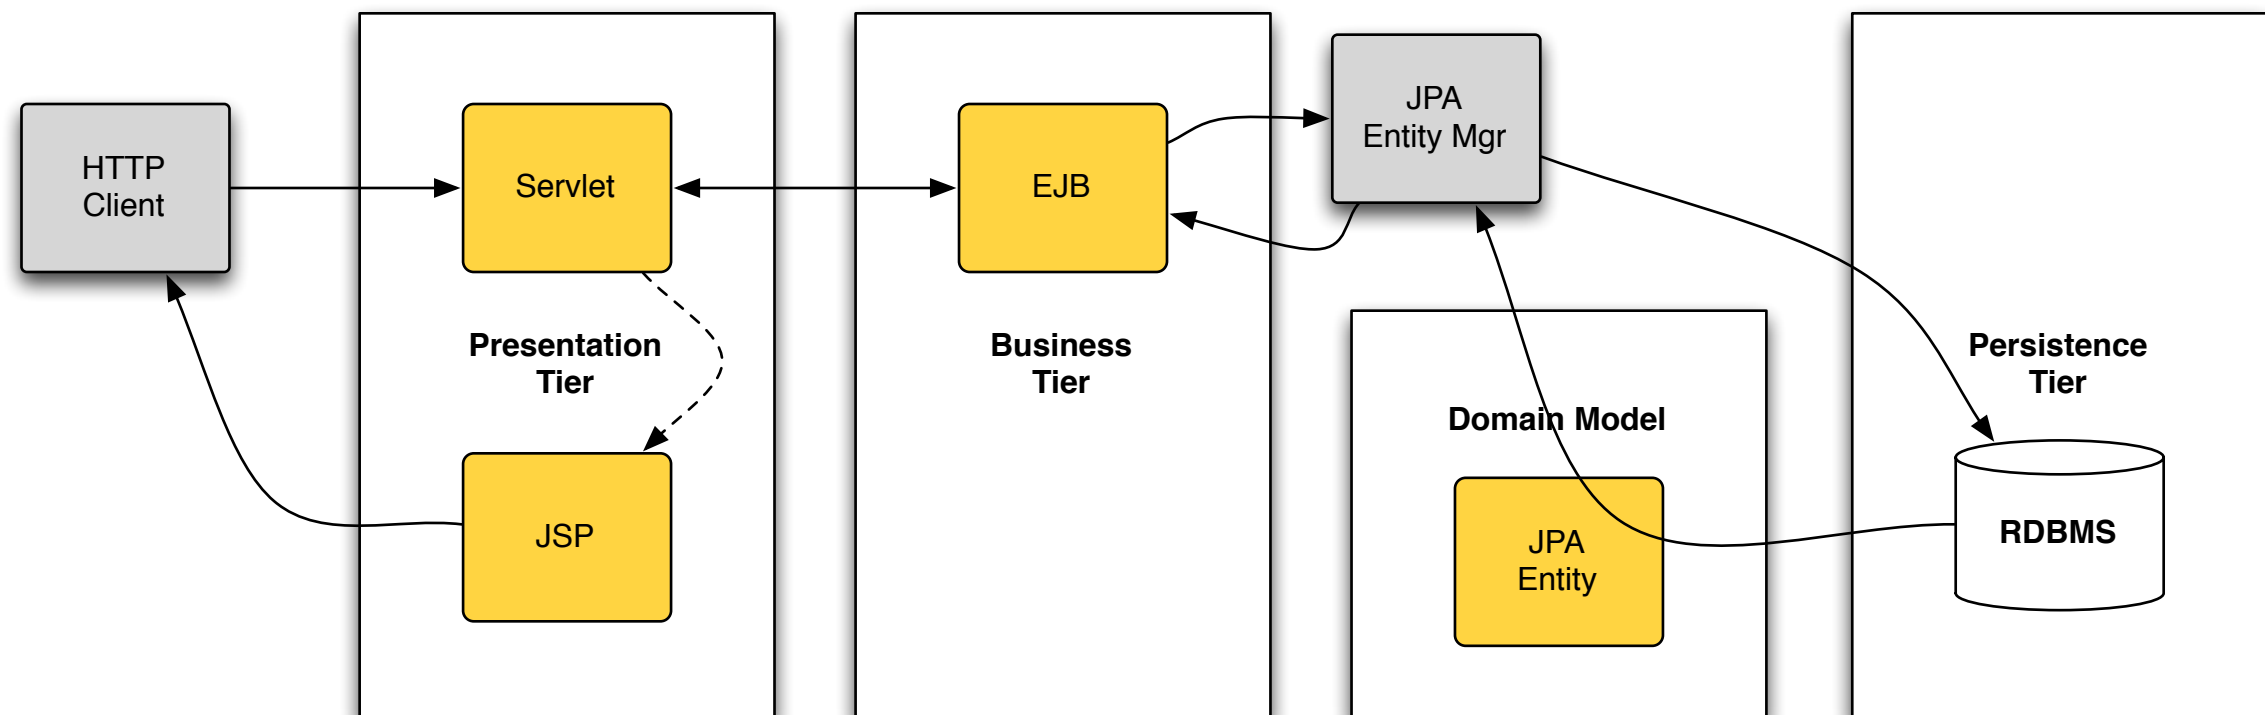
    ...
}
```



Reference Architecture



Reference Architecture



Reference Architecture

